

# Model-Free Reinforcement Learning

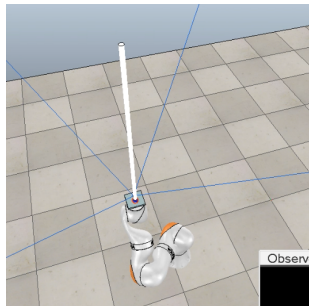

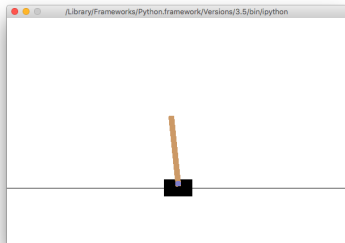
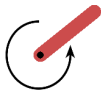


Mathias Winther Madsen  
`mathias@micropsi-industries.com`

March 6, 2017

# Policy Gradient Methods

---



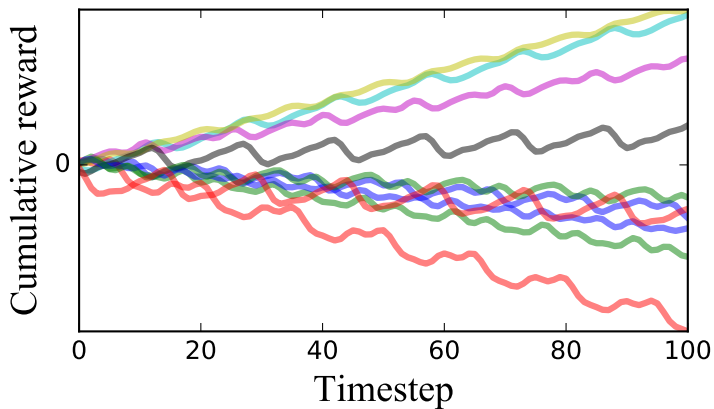
# Policy Gradient Methods

The REINFORCE algorithm: Williams, “Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning” (*Machine Learning*, 1992)

| Task                      | Random        | REINFORCE            | TNPG                  | RWR                  | REPS           | TRPO                  | CEM                | CMA-ES             | DDPG                  |
|---------------------------|---------------|----------------------|-----------------------|----------------------|----------------|-----------------------|--------------------|--------------------|-----------------------|
| Cart-Pole Balancing       | 77.1 ± 0.0    | 4693.7 ± 14.0        | <b>3986.4 ± 748.9</b> | <b>4861.5 ± 12.3</b> | 565.6 ± 137.6  | <b>4869.8 ± 37.6</b>  | 4815.4 ± 4.8       | 2440.4 ± 568.3     | 4634.4 ± 87.8         |
| Inverted Pendulum*        | -153.4 ± 0.2  | 13.4 ± 18.0          | <b>209.7 ± 55.5</b>   | 84.7 ± 13.8          | -113.3 ± 4.6   | <b>247.2 ± 76.1</b>   | 38.2 ± 25.7        | -40.1 ± 5.7        | 40.0 ± 244.6          |
| Mountain Car              | -415.4 ± 0.0  | -67.1 ± 1.0          | <b>-66.5 ± 4.5</b>    | -79.4 ± 1.1          | -275.6 ± 166.3 | <b>-61.7 ± 0.9</b>    | -66.0 ± 2.4        | -85.0 ± 7.7        | -288.4 ± 170.3        |
| Acrobot                   | -1904.5 ± 1.0 | -508.1 ± 91.0        | -395.8 ± 121.2        | -352.7 ± 35.9        | -1001.5 ± 10.8 | -326.0 ± 24.4         | -436.8 ± 14.7      | -785.6 ± 13.1      | <b>-233.6 ± 5.8</b>   |
| Double Inverted Pendulum* | 149.7 ± 0.1   | 4116.5 ± 65.2        | <b>4455.4 ± 37.6</b>  | 3614.8 ± 368.1       | 446.7 ± 114.8  | <b>4412.4 ± 50.4</b>  | 2566.2 ± 178.9     | 1576.1 ± 51.3      | 2863.4 ± 154.0        |
| Swimmer*                  | -1.7 ± 0.1    | 92.3 ± 0.1           | <b>96.0 ± 0.2</b>     | 60.7 ± 5.5           | 3.8 ± 3.3      | <b>96.0 ± 0.2</b>     | 68.8 ± 2.4         | 64.9 ± 1.4         | 85.8 ± 1.8            |
| Hopper                    | 8.4 ± 0.0     | 714.0 ± 29.3         | <b>1155.1 ± 57.9</b>  | 553.2 ± 71.0         | 86.7 ± 17.6    | <b>1183.3 ± 150.0</b> | 63.1 ± 7.8         | 20.3 ± 14.3        | 267.1 ± 43.5          |
| 2D Walker                 | -1.7 ± 0.0    | 506.5 ± 78.8         | <b>1382.6 ± 108.2</b> | 136.0 ± 15.9         | -37.0 ± 38.1   | <b>1353.8 ± 85.0</b>  | 84.5 ± 19.2        | 77.1 ± 24.3        | 318.4 ± 181.6         |
| Half-Cheetah              | -90.8 ± 0.3   | 1183.1 ± 69.2        | <b>1729.5 ± 184.6</b> | 376.1 ± 28.2         | 34.5 ± 38.0    | <b>1914.0 ± 120.1</b> | 330.4 ± 274.8      | 441.3 ± 107.6      | <b>2148.6 ± 702.7</b> |
| Ant*                      | 13.4 ± 0.7    | 548.3 ± 55.5         | <b>706.0 ± 127.7</b>  | 37.6 ± 3.1           | 39.0 ± 9.8     | <b>730.2 ± 61.3</b>   | 42.2 ± 5.9         | 17.8 ± 15.5        | 326.2 ± 20.8          |
| Simple Humanoid           | 41.5 ± 0.2    | 128.1 ± 34.0         | <b>255.0 ± 24.5</b>   | 93.3 ± 17.4          | 28.3 ± 4.7     | <b>269.7 ± 40.3</b>   | 60.6 ± 12.9        | 28.7 ± 3.9         | 99.4 ± 28.1           |
| Full Humanoid             | 13.2 ± 0.1    | 262.2 ± 10.5         | <b>288.4 ± 25.2</b>   | 46.7 ± 5.6           | 41.7 ± 6.1     | <b>287.0 ± 23.4</b>   | 36.9 ± 2.9         | N/A ± N/A          | 119.0 ± 31.2          |
| Cart-Pole Balancing (LS)* | 77.1 ± 0.0    | 420.9 ± 265.5        | <b>945.1 ± 27.8</b>   | 68.9 ± 1.5           | 898.1 ± 22.1   | <b>960.2 ± 46.0</b>   | 227.0 ± 223.0      | 68.0 ± 1.6         |                       |
| Inverted Pendulum (LS)    | -122.1 ± 0.1  | -13.4 ± 3.2          | <b>0.7 ± 6.1</b>      | -107.4 ± 0.2         | -87.2 ± 8.0    | <b>4.5 ± 4.1</b>      | -81.2 ± 33.2       | -62.4 ± 3.4        |                       |
| Mountain Car (LS)         | -83.0 ± 0.0   | -81.2 ± 0.6          | <b>-65.7 ± 9.0</b>    | -81.7 ± 0.1          | -82.6 ± 0.4    | <b>-64.2 ± 9.5</b>    | <b>-68.9 ± 1.3</b> | <b>-73.2 ± 0.6</b> |                       |
| Acrobot (LS)*             | -393.2 ± 0.0  | -128.9 ± 11.6        | <b>-84.6 ± 2.9</b>    | -235.9 ± 5.3         | -379.5 ± 1.4   | <b>-83.3 ± 9.9</b>    | -149.5 ± 15.3      | -159.9 ± 7.5       |                       |
| Cart-Pole Balancing (NO)* | 101.4 ± 0.1   | 616.0 ± 210.8        | <b>916.3 ± 23.0</b>   | 93.8 ± 1.2           | 99.6 ± 7.2     | 606.2 ± 122.2         | 181.4 ± 32.1       | 104.4 ± 16.0       |                       |
| Inverted Pendulum (NO)    | -122.2 ± 0.1  | 6.5 ± 1.1            | <b>11.5 ± 0.5</b>     | -110.0 ± 1.4         | -119.3 ± 4.2   | <b>10.4 ± 2.2</b>     | -55.6 ± 16.7       | -80.3 ± 2.8        |                       |
| Mountain Car (NO)         | -83.0 ± 0.0   | -74.7 ± 7.8          | <b>-64.5 ± 8.6</b>    | -81.7 ± 0.1          | -82.9 ± 0.1    | <b>-60.2 ± 2.0</b>    | -67.4 ± 1.4        | -73.5 ± 0.5        |                       |
| Acrobot (NO)*             | -393.5 ± 0.0  | <b>-186.7 ± 31.3</b> | <b>-164.5 ± 13.4</b>  | -233.1 ± 0.4         | -258.5 ± 14.0  | <b>-149.6 ± 8.6</b>   | -213.4 ± 6.3       | -236.6 ± 6.2       |                       |
| Cart-Pole Balancing (SI)* | 76.3 ± 0.1    | 431.7 ± 274.1        | <b>980.5 ± 7.3</b>    | 69.0 ± 2.8           | 702.4 ± 196.4  | <b>980.3 ± 5.1</b>    | 746.6 ± 93.2       | 71.6 ± 2.9         |                       |
| Inverted Pendulum (SI)    | -121.8 ± 0.2  | -5.3 ± 5.6           | <b>14.8 ± 1.7</b>     | -108.7 ± 4.7         | -92.8 ± 23.9   | <b>14.1 ± 0.9</b>     | -51.8 ± 10.6       | -63.1 ± 4.8        |                       |
| Mountain Car (SI)         | -82.7 ± 0.0   | -63.9 ± 0.2          | <b>-61.8 ± 0.4</b>    | -81.4 ± 0.1          | -80.7 ± 2.3    | <b>-61.6 ± 0.4</b>    | -63.9 ± 1.0        | -66.9 ± 0.6        |                       |
| Acrobot (SI)*             | -387.8 ± 1.0  | <b>-169.1 ± 32.3</b> | <b>-156.6 ± 38.9</b>  | -233.2 ± 2.6         | -216.1 ± 7.7   | <b>-170.9 ± 40.3</b>  | -250.2 ± 13.7      | -245.0 ± 5.5       |                       |
| Swimmer + Gathering       | 0.0 ± 0.0     | 0.0 ± 0.0            | 0.0 ± 0.0             | 0.0 ± 0.0            | 0.0 ± 0.0      | 0.0 ± 0.0             | 0.0 ± 0.0          | 0.0 ± 0.0          | 0.0 ± 0.0             |
| Ant + Gathering           | -5.8 ± 5.0    | -0.1 ± 0.1           | -0.4 ± 0.1            | -5.5 ± 0.5           | -6.7 ± 0.7     | -0.4 ± 0.0            | -4.7 ± 0.7         | N/A ± N/A          | -0.3 ± 0.3            |
| Swimmer + Maze            | 0.0 ± 0.0     | 0.0 ± 0.0            | 0.0 ± 0.0             | 0.0 ± 0.0            | 0.0 ± 0.0      | 0.0 ± 0.0             | 0.0 ± 0.0          | 0.0 ± 0.0          | 0.0 ± 0.0             |
| Ant + Maze                | 0.0 ± 0.0     | 0.0 ± 0.0            | 0.0 ± 0.0             | 0.0 ± 0.0            | 0.0 ± 0.0      | 0.0 ± 0.0             | 0.0 ± 0.0          | N/A ± N/A          | 0.0 ± 0.0             |

Table: Duan, Chen, Houthoofd, Schulman and Abbeel, “Benchmarking Deep Reinforcement Learning for Continuous Control,” *Proceedings of the ICML*, 2016.


## Policy Gradient Methods



# One-Shot Games

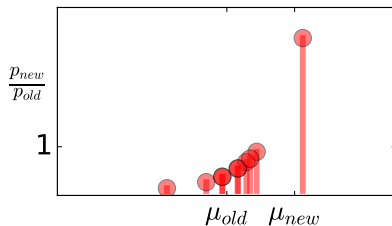
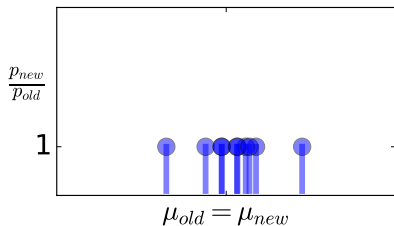
Problem statement:

$$W^* = \arg \max_W E_W[V]$$



# Importance Weighting

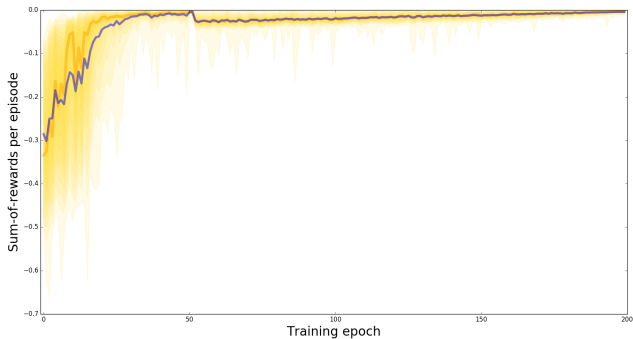
$$E_{new}[X] = E_{old}\left[X \cdot \frac{p_{new}}{p_{old}}\right]$$



# The Policy Gradient

$$\nabla_W E_W[V] = E_{W_0} \left[ V \cdot \frac{\nabla_W p_W}{p_{W_0}} \right]$$


Dart-Throwing Game:  $R(u) = -\|u - u^*\|^2$





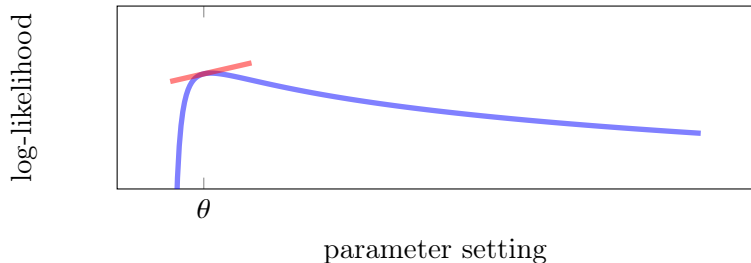
# The Policy Gradient

$$\nabla_{\mathbf{W}} E_{\mathbf{W}} [V] = E_{\mathbf{W}_0} \left[ V \cdot \left( \frac{\nabla_{\mathbf{W}} P}{P} \right) \right]$$



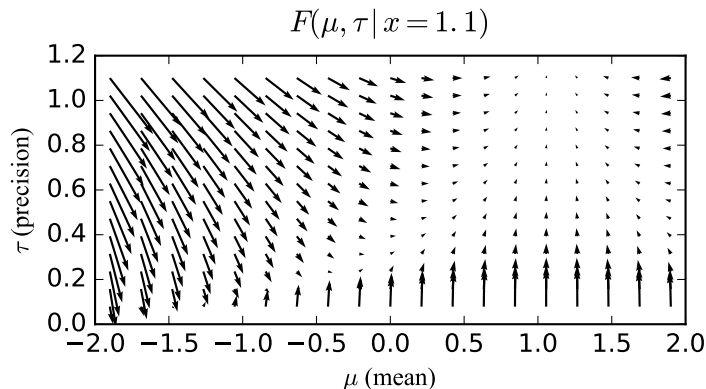
# The Fisher Score

$$F(x) = \frac{\nabla_{\theta} d(x | \theta)}{d(x | \theta)} = \nabla_{\theta} \log d(x | \theta)$$




# The Fisher Score


$$\nabla_{(\mu, \tau)} \log \left( \sqrt{\frac{\tau}{\pi}} \exp \left\{ -\tau(x - \mu)^2 \right\} \right) = \begin{pmatrix} 2\tau(x - \mu) \\ (2\tau)^{-1} - (x - \mu)^2 \end{pmatrix}$$



# Extensive Games



# Extensive Games



$\text{Action} \sim \text{Distribution}(\text{History}, \text{Weights})$

# Extensive Games


Problem statement:

$$W^* = \arg \max_w E_w [R_0 + R_1 + \cdots + R_{T-1}]$$

# The Policy Gradient

$$\nabla_W E_W[V] = E_W[V \cdot F]$$

## Episode Scores



$$q(S_0) p(U_0 | S_0, W) q(S_1 | S_0, U_0) p(U_1 | S_1, W) q(S_2 | S_1, U_1) \dots$$



## Episode Scores

$$\frac{\nabla (q_0 p_1 q_1 p_1 q_2 p_2 \cdots q_{T-1} p_{T-1})}{(q_0 p_1 q_1 p_1 q_2 p_2 \cdots q_{T-1} p_{T-1})} = \frac{\nabla (p_1 p_1 p_2 \cdots p_{T-1})}{(p_1 p_1 p_2 \cdots p_{T-1})}$$

Hence:

$$F = \nabla \log p_0 + \nabla \log p_1 + \nabla \log p_2 + \cdots + \nabla \log p_{T-1}$$

# The Policy Gradient

$$\nabla_W E_W \left[ \sum_{t=0}^{T-1} R_t \right] = E_W \left[ \underbrace{\sum_{t=0}^{T-1} R_t}_V \cdot \underbrace{\sum_{t=0}^{T-1} \nabla_W \log p(U_t | S_t, W)}_F \right]$$


# Repeat-After-Me Game: $R(s, u) = -\|s - u\|^2$

$$s \sim \mathcal{N}(1/2, 1)$$


$$u \sim \mathcal{N}(as + b, 1)$$

$$F\left(\begin{matrix} a \\ b \end{matrix} \middle| s, u\right) = \begin{pmatrix} (as + b - u)s \\ (as + b - u) \end{pmatrix}$$


Score estimates



Gradient estimates




# Repeat-After-Me Game: $R(s, u) = -\|s - u\|^2$



# Apportioning Blame

$$V = R_0 + R_1 + R_2 + \cdots + R_{T-1}$$

$$F = F_0 + F_1 + F_2 + \cdots + F_{T-1}$$




# Apportioning Blame

$$E \begin{bmatrix} F_0 R_0 & F_0 R_1 & F_0 R_2 & \cdots & F_0 R_{T-1} \\ F_1 R_0 & F_1 R_1 & F_1 R_2 & \cdots & F_1 R_{T-1} \\ F_2 R_0 & F_2 R_1 & F_2 R_2 & \cdots & F_2 R_{T-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ F_{T-1} R_0 & F_{T-1} R_1 & F_{T-1} R_2 & \cdots & F_{T-1} R_{T-1} \end{bmatrix}$$

# Apportioning Blame

$$E \begin{bmatrix} F_0 R_0 & F_0 R_1 & F_0 R_2 & \cdots & F_0 R_{T-1} \\ 0 & F_1 R_1 & F_1 R_2 & \cdots & F_1 R_{T-1} \\ 0 & 0 & F_2 R_2 & \cdots & F_2 R_{T-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & F_{T-1} R_{T-1} \end{bmatrix}$$

# Apportioning Blame



Keeping or dropping  
the zero-mean terms.



# The Policy Gradient Method

For  $I$  epochs:

For  $N$  episodes:

Collect rollout,  $[(S_t, U_t, R_t)]_{t=0}^{T-1}$

For each action  $U_t$ :

Compute score,  $F_t = F(W | U_t)$ ;

Compute tailsum,  $V_t = \sum_{k=t}^{T-1} R_k$ ;

Get gradient,  $G_n = \sum_{t=0}^{T-1} V_t F_t$ ;

Average gradients,  $G = \frac{1}{N} \sum_{n=0}^{N-1} G_n$ ;

Adjust weights,  $W \leftarrow W + \alpha G$ ;

# Future Complications

1. Short-lived dependencies:

$$V_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \gamma^3 R_{t+3} + \cdots$$

2. Baselines:

$$V_t \leftarrow V_t - \hat{V}_t$$

3. “Natural” directions:

$$G \leftarrow E[-\nabla F]^{-1} \cdot G$$

4. Line searches:

$$\alpha = \arg \max_{\alpha} J(W + \alpha G)$$

# References

- ▶ Fisher: “On the Mathematical Foundations of Theoretical Statistics” (*Proceedings of the Royal Society A*, 1922).
- ▶ Williams: “Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning” (*Machine Learning*, 1992).
- ▶ Amari: “Natural Gradient Works Efficiently in Learning” (*Neural Computation*, 1998).
- ▶ Schulman, Levine, Moritz, Jordan, and Abbeel: “Trust Region Policy Optimization” (*ICML*, 2015).
- ▶ Duan, Chen, Houthoofd, Schulman, and Abbeel, “Benchmarking Deep Reinforcement Learning for Continuous Control” (*ICML*, 2016).

`github.com/mathias-madsen/reinforce_tutorial`